

1 Statistical Issues in Multicenter Randomized Clinical
2 Trials

3 Ronald G. Thomas, Ph.D.*

4 August 14, 2026

5 **Abstract**

6 Multicenter randomized controlled trials require a choice among ignor-
7 ing, fixing, or randomizing site effects in the primary analysis, and this
8 choice interacts with stratified-randomization adjustment, treatment-by-site
9 interaction, and site-size balance. We report a factorial Monte Carlo sim-
10 ulation, structured under the ADEMP framework of Morris, White, and
11 Crowther (2019), that jointly varies the number of sites, site-level variance,
12 treatment-by-site interaction, and site-size balance, comparing OLS ignoring
13 site, OLS with fixed site effects, and a linear mixed model with random site
14 effects. In addition to bias, empirical and model-based standard error, power,
15 and coverage under a true alternative treatment effect, we report empirical
16 Type I error under a true null effect, and singular/boundary-fit rates for
17 the random-effects model, both previously unreported in this simulation de-
18 sign. The random-effects model attains the best combination of power and
19 coverage across most conditions; ignoring site produces conservative (not
20 anti-conservative) inference under stratified randomization, consistent with
21 prior work; and all three methods control Type I error near the nominal
22 level in the null-hypothesis scenarios simulated. For the treatment-by-site-
23 interaction scenarios, where the random-intercept model is misspecified, we
24 additionally fit and compare a random-slopes model, finding a modest cover-
25 age change but a substantially higher singular-fit rate, a trade-off not previ-
26 ously quantified in this simulation design. Every point estimate is reported
27 with its Morris et al. Table 6 Monte Carlo standard error. Code and a full
28 ADEMP audit trail accompany the simulation.

29 **Contents**

30	1 Introduction	2
31	1.1 Modeling site effects	2
32	1.2 What does “random” actually mean here?	3
33	1.3 Failure to adjust for stratification variables	5
34	1.4 Treatment-by-site interaction	6
35	1.5 Imbalanced site sizes	6
36	1.6 Present study	7

*Wertheim School of Public Health, UCSD; ORCID: 0000-0003-1686-4965

37	2 Methods	7
38	2.1 ADEMP structure	7
39	2.2 Data generating model	8
40	2.3 Parameter values	8
41	2.4 Analytic methods	8
42	2.5 Simulation procedure	9
43	3 Results	9
44	3.1 Headline findings	9
45	3.2 Summary of performance metrics	10
46	3.3 Bias	10
47	3.4 Power and coverage	10
48	3.5 Impact of imbalanced site sizes	10
49	3.6 Type I error under the null	10
50	3.7 Random slopes under treatment-by-site interaction	13
51	4 Discussion	14
52	5 Future research	16
53	6 Data and code availability	17
54	7 Conflict of interest	17
55	8 References	17
56	8.1 Morris et al. (2019) ADEMP Compliance	17

57 1 Introduction

58 Multicenter randomized clinical trials (RCTs) represent the standard design for
59 confirmatory evaluation of therapeutic interventions. By enrolling participants
60 across multiple clinical sites, these trials achieve practical recruitment targets
61 while broadening the population base for generalization of findings^{1,2}. However,
62 the multicenter structure introduces several analytic complexities that, if mishan-
63 dled, can distort inference about treatment effects. We address four interrelated
64 statistical issues that arise in the design and analysis of multicenter RCTs with
65 continuous outcomes, and we present simulation evidence bearing on each in turn.

66 1.1 Modeling site effects

67 The first question concerns how site (or center) effects should be incorporated
68 into the primary analysis. Three broad strategies exist: ignoring site differences
69 entirely, treating sites as fixed effects, and treating sites as random effects drawn
70 from a population of potential sites. Each strategy carries distinct assumptions
71 and implications for the validity and efficiency of treatment effect estimation^{3,4}.

72 Localio and colleagues provided an influential overview of the problem, identifying
73 three analytic challenges inherent to multicenter data: within-center correlation

of outcomes, confounding by center, and effect modification across centers². The choice between fixed and random effects models has been debated extensively. Fixed effects models condition on the observed sites and make no distributional assumptions about site effects, but they consume degrees of freedom rapidly when the number of sites is large relative to total sample size⁵. Random effects models treat site effects as realizations from a normal distribution, estimating only a variance component rather than individual site parameters, and thereby preserve statistical efficiency⁶. The ICH E9 guideline recommends that mixed models are ‘especially relevant when the number of sites is large’¹.

Chu and colleagues conducted a comprehensive simulation comparing six analytic approaches for continuous outcomes across varying numbers of centers, center sizes, and intraclass correlation coefficients (ICCs)⁷. All methods yielded unbiased point estimates, but ignoring centers inflated the standard error and reduced power when within-center correlation was present. The mixed-effects model attained nominal coverage and near-optimal power across nearly all configurations. Generalized estimating equations (GEE) underestimated the standard error when the number of centers was small.

Kahan and Morris extended these findings, demonstrating that random center effects performed at least as well as fixed effects in all scenarios examined, and substantially outperformed them when the number of patients per center was small⁸. Islam and Bangdiwala confirmed these patterns in a recent replication study with both continuous and binary outcomes, reinforcing the recommendation for random effects models as a default analytic strategy⁹.

1.2 What does “random” actually mean here?

Before going further, it is worth being precise about the word “random,” because this word carries most of the weight in the discussion above, and because it can mean two different things. We spell out both meanings in plain terms, with the underlying math shown alongside, since the two meanings lead to two different arguments later in this Introduction.

1.2.1 A hospital-specific shift, fixed or random

Picture a trial run at ten hospitals. Each hospital has its own typical outcome level, for reasons that have nothing to do with the treatment being tested: a different mix of patients, different routine practices, different equipment. Call that hospital-specific shift α_i , for hospital i . A simple model for one patient’s outcome, written Y_{ij} for patient j at hospital i , looks like this:

$$Y_{ij} = \alpha_i + \delta T_{ij} + \epsilon_{ij}$$

Here T_{ij} is 1 if that patient received the active treatment and 0 if not, δ is the treatment effect the trial wants to estimate, and ϵ_{ij} is ordinary patient-to-patient variation that has nothing to do with which hospital a patient attended.

112 The open question is what kind of quantity α_i is.

113 There are two options. The first treats each hospital's shift as an ordinary, fixed
114 number, no different in kind from a patient's age or sex: something to adjust for,
115 one number per hospital. Ten hospitals means ten fixed numbers to estimate.

116 The second option treats α_i as random, drawn from a normal distribution: $\alpha_i \sim$
117 $N(0, \sigma_\alpha^2)$. Under this option, the one unknown quantity of real interest is σ_α^2 : how
118 much do hospitals typically differ from one another? Rather than estimating ten
119 separate numbers, the analysis estimates a single number describing the *spread* of
120 hospital effects.

121 This is not a small technical difference. It changes what the analysis is entitled
122 to claim. Estimating ten fixed numbers tells you only about these ten specific
123 hospitals. Estimating one spread, σ_α^2 , is a broader claim: it says something about
124 hospital-to-hospital variation in general, of the kind you would expect if the trial
125 had, by chance, enrolled a different set of ten hospitals instead.

126 That broader claim is exactly where a disagreement in the published literature
127 shows up, and we come back to it below.

128 1.2.2 Why the simulation needs a genuinely random site effect

129 This paper is a computer simulation. It generates many pretend trials, over and
130 over, and checks how well each analysis strategy (ignore site, treat site as fixed,
131 treat site as random) recovers the true treatment effect and reports honest uncer-
132 tainty around it. For that check to mean anything, the computer program that
133 generates each pretend trial has to actually build in hospital-to-hospital variation,
134 and it has to do so afresh for every pretend trial, not use the same ten fixed
135 numbers every time.

136 That is what drawing $\alpha_i \sim N(0, \sigma_\alpha^2)$ inside the simulation accomplishes: every
137 pretend trial gets its own new set of hospital shifts, all coming from a distribution
138 with the same spread, σ_α^2 . Only by rebuilding this variation from scratch each
139 time can the simulation ask a fair question: across many possible versions of “a
140 multicenter trial like this one,” does each analysis strategy give honest confidence
141 intervals and adequate power? A simulation that reused the same fixed hospital
142 numbers every time would not be testing that question at all; it would only be
143 testing how well each method recovers one specific, already-fixed set of numbers.
144 The later sections on treatment-by-site interaction and on imbalanced site sizes
145 work the same way: each also needs a freshly drawn random ingredient, every
146 pretend trial, for the same reason.

147 It is important to be clear about what this operational choice does, and does
148 not, claim. It is a requirement of building a fair test, not a statement about how
149 real hospitals were chosen for a real trial. That second, separate question is the
150 subject of the next section.

1.2.3 Two different reasons statisticians give for treating site as random

Everything so far concerns how the simulation is built. A separate question is why, in a real trial with real hospitals, an analyst should choose to treat site as random rather than fixed when fitting the actual statistical model. Published statisticians do not fully agree on the reason, even where they agree on the recommendation, and a statistically literate reader deserves to know both arguments.

The first argument says the random-effects model literally means what it appears to mean: it treats the hospitals in the trial as if they had been drawn at random from a much larger population of possible hospitals. Falissard and Chavance make this point directly: the random-effects interpretation is only truly justified when the hospitals really were chosen at random from a wider population¹⁰. In practice, that is almost never how hospitals are chosen. Hospitals join a trial because they can recruit patients quickly, because an investigator has a relationship with the study team, or because of funding and logistics, not by a lottery over all possible hospitals³.

The second argument does not depend on how the hospitals were chosen at all. Kahan and Morris justify adjusting for site on the design of the trial itself: patients were randomly split between treatment and control separately within each hospital, so the analysis should reflect that split, no matter how the hospitals themselves came to be in the trial^{8,11}. Edgar, Roberts, and Sharples reach the same conclusion by re-analyzing a real trial run across 271 hospitals in 40 countries: they recommend a random site effect for practical reasons, better power and less bias, again without any claim about hospitals being a random sample of a larger population¹².

These two arguments do not agree on *why* a random site effect is the right choice, but they agree on *what to do*: prefer a random effect for site once there are more than a handful of hospitals. The recommendation in this paper rests on that agreement, so it stands regardless of which argument a reader finds more convincing. The simulation itself follows the second, design-based tradition in spirit: it builds hospital-to-hospital variation into every pretend trial as an operational necessity for testing the analysis strategies fairly (see above), without needing or claiming that real hospitals are ever a literal random sample of all possible hospitals.

1.3 Failure to adjust for stratification variables

The second issue concerns the consequences of ignoring stratification variables, including site, in the analysis. When randomization is stratified by center (as is standard in multicenter trials), failing to adjust for the stratification factor in the analysis leads to biased inference. Kahan and Morris demonstrated through simulation that an unadjusted analysis of a trial randomized with stratified blocks produces standard errors that are biased upward, confidence intervals that are too wide, type I error rates below the nominal level, and a loss of statistical power¹¹. This paradoxical conservatism arises because the stratified randomization induces

negative correlation between treatment groups within strata; ignoring this correlation inflates variance estimates.

Kahan and Morris extended this work to trials with multiple stratification factors, finding that adjusting for the main effects of all stratification variables was generally sufficient without requiring adjustment for all cross-classified strata¹³. These findings align with the ICH E9 recommendation that ‘the statistical model used for analysis should reflect the restriction imposed by the randomisation procedure’¹, and with current regulatory guidance on covariate adjustment more broadly¹⁴.

1.4 Treatment-by-site interaction

The third concern is heterogeneity of the treatment effect across sites. Even under a common protocol, eligible participants, site practices, and investigator experience may vary, producing genuine treatment-by-site interaction¹⁵. Senn argued that some degree of treatment-by-center interaction is inevitable in practice and that a rational approach to planning multicenter trials should anticipate it³.

The ICH E9 guideline recommends first fitting a model without the interaction term to estimate the main treatment effect, and then examining heterogeneity as a secondary analysis. However, the power to detect treatment-by-site interaction is typically low, so a non-significant interaction test does not establish homogeneity¹⁶. Jones and colleagues compared fixed-effects weighted estimators, unweighted estimators, and random effects estimators in the presence of varying degrees of treatment-by-center interaction⁴. The fixed effects weighted estimator performed well across a range of interaction magnitudes, while unweighted estimators could be inefficient when site sizes differed markedly.

When treatment-by-site interaction is modeled as random (i.e., the treatment slope varies across sites), the mixed model estimates both the average treatment effect and the variance of treatment effects across sites, providing a natural summary of heterogeneity¹⁷.

1.5 Imbalanced site sizes

The fourth issue, closely related to the preceding three, is whether imbalance in site sizes affects inference. Multicenter trials frequently exhibit highly unequal enrollment across sites, with a few large sites contributing a disproportionate share of participants. Senn noted that a rational approach to planning multicenter trials leads naturally to unequal site sizes, because sites with higher recruitment capacity will enroll more patients within the trial’s fixed time window³.

Ruvuna introduced a coefficient of imbalance to quantify the power loss attributable to unequal center sizes and showed that substantial imbalance can reduce power when an unweighted (Type III) analysis is used¹⁸. Vierron and Giraudeau formalized this through a design effect framework, demonstrating that the efficiency loss depends on both the ICC and a statistic quantifying the heterogeneity of group distributions across centers¹⁹. When the ICC is positive, balanced allocation is optimal, but moderate imbalance imposes only a modest

234 efficiency penalty. The key risk emerges when site size is informative (that is,
235 when larger sites systematically differ from smaller sites in patient characteristics
236 or treatment delivery), potentially introducing bias into the treatment effect
237 estimate^{15,20}.

238 1.6 Present study

239 The existing literature provides clear theoretical guidance but limited integrated
240 simulation evidence spanning all four issues simultaneously. We present a Monte
241 Carlo simulation that jointly varies the analytic method (ignore site, fixed site
242 effects, random site effects), the presence or absence of treatment-by-site inter-
243 action, balanced versus imbalanced site sizes, and the magnitude of site-level
244 variance. The simulation evaluates bias, empirical standard error, model-based
245 standard error, power, and coverage probability of 95% confidence intervals for the
246 average treatment effect across all factorial combinations of these design features.

247 2 Methods

248 2.1 ADEMP structure

249 Following Morris, White, and Crowther²¹, the simulation is reported under the
250 ADEMP framework.

- 251 • **Aims.** Compare three primary analytic strategies for handling site het-
252 erogeneity in multicenter RCTs (ignore site, fixed site effects, random site
253 effects) across scenarios varying the number of sites, site variance, treatment-
254 by-site interaction, and site-size balance; separately assess whether these
255 strategies control the Type I error rate under a true null treatment effect;
256 and, for the treatment-by-site-interaction scenarios specifically, compare the
257 random-intercept specification against a fourth, correctly specified random-
258 slopes method.
- 259 • **Data-generating mechanism.** Detailed in the next subsection.
- 260 • **Estimand.** The overall treatment effect δ on the outcome scale, in the
261 simulation-performance sense of Morris et al. (a data-generating parameter
262 with a known true value), not in the ICH E9(R1) sense of a fully specified
263 population/variable/ intercurrent-event/summary-measure estimand²².
- 264 • **Methods.** OLS ignoring site, OLS with fixed site effects, a linear mixed
265 model with random site effects, and (fit on every replicate for a paired com-
266 parison, but reported only for the treatment-by-site-interaction scenarios) a
267 linear mixed model with random site effects and a random treatment slope,
268 all via `lme4::lmer`. All four methods are tested and interval-estimated us-
269 ing a *t*-reference distribution: the OLS residual degrees of freedom for the
270 ignore-site and fixed-site methods, and *N* minus the number of fixed-effect
271 parameters for the two mixed-model methods. The latter is a naive approx-
272 imation, not a Kenward-Roger or Satterthwaite denominator-df correction
273 (see Future research); it is used consistently for both the *p*-value and the
274 95% CI half-width for both mixed-model methods, replacing an earlier im-

plementation in which the p -value used this t -approximation but the CI used a fixed $z = 1.96$ half-width regardless of sample size or method.

- **Performance measures.** Bias, empirical SE, MSE, mean model SE, power at $\alpha = 0.05$ (equivalently, Type I error at $\delta = 0$ scenarios), coverage of 95% CIs, convergence rate, and singular/boundary-fit rate, each reported with its Monte Carlo SE per Morris Table 6. Non-convergence (a hard error from `lmer`) and singular/boundary fits (a successful fit at the edge of the random-effects covariance parameter space, which `lme4` flags as a warning rather than an error) are tracked as two distinct first-class outcomes and reported per scenario and method, rather than treating a fit that raised any warning as simply non-convergent or, conversely, absorbing a boundary fit silently into “converged” with no record that it was a boundary solution.

2.2 Data generating model

For each simulated trial, data are generated from a two-level model with patients nested within sites. Let Y_{ij} denote the continuous outcome for patient j at site i :

$$Y_{ij} = \alpha_i + (\delta + \gamma_i)T_{ij} + \epsilon_{ij}$$

where $\alpha_i \sim N(0, \sigma_\alpha^2)$ is the random site intercept, $\delta = 0.30$ is the true average treatment effect, $\gamma_i \sim N(0, \sigma_\gamma^2)$ is the random treatment-by-site interaction, $T_{ij} \in \{0, 1\}$ is the treatment indicator (balanced within each site via permuted blocks), and $\epsilon_{ij} \sim N(0, \sigma_e^2 = 1)$ is the residual error.

2.3 Parameter values

The simulation crosses the following factors:

- **Number of sites:** $K = 10$ (few) vs $K = 30$ (many)
- **Site-level variance:** $\sigma_\alpha^2 \in \{0, 0.10, 0.50\}$ (none, moderate, large)
- **Treatment-by-site interaction:** $\sigma_\gamma^2 \in \{0, 0.04\}$ (none vs present)
- **Site size balance:** balanced ($n_i = 20$ for all i) vs imbalanced (site sizes drawn from a $\text{Gamma}(0.8, 0.8/20)$ distribution, yielding a coefficient of variation ≈ 1.1)

The true treatment effect is $\delta = 0.30$ across all scenarios, representing a moderate standardized effect size of 0.30 standard deviations.

2.4 Analytic methods

Each simulated dataset is analyzed by three primary methods, plus a fourth fit on every replicate for a paired comparison but reported only for the treatment-by-site-interaction scenarios (Results, “Random slopes under treatment-by-site interaction”):

1. **Ignore site:** $Y_{ij} = \beta_0 + \beta_1 T_{ij} + e_{ij}$ (ordinary least squares, no site term)

2. **Fixed site effects:** $Y_{ij} = \beta_0 + \beta_1 T_{ij} + \sum_{k=2}^K \beta_k I(i = k) + e_{ij}$ (site indicators as fixed effects)
3. **Random site effects:** $Y_{ij} = \beta_0 + \beta_1 T_{ij} + u_i + e_{ij}$ where $u_i \sim N(0, \sigma_u^2)$ (linear mixed model via REML, fit with `lme4`)
4. **Random slopes:** $Y_{ij} = \beta_0 + \beta_1 T_{ij} + u_i + v_i T_{ij} + e_{ij}$ where (u_i, v_i) jointly follow a bivariate normal distribution (random intercept and random treatment slope per site, via REML). This is the correctly specified model under the treatment-by-site-interaction data-generating mechanism (Data generating model, γ_i), unlike Method 3, which omits the random slope and is misspecified whenever $\sigma_\gamma^2 > 0$.

2.5 Simulation procedure

Scenario 1/24: K10_noVar_bal ... Scenario 2/24: K10_modVar_bal ... Scenario 3/24: K10_hiVar_bal ... Scenario 4/24: K10_modVar_imbal ... Scenario 5/24: K10_hiVar_imbal ... Scenario 6/24: K10_modVar_int ... Scenario 7/24: K10_hiVar_int ... Scenario 8/24: K10_hiVar_int_imb ... Scenario 9/24: K30_noVar_bal ... Scenario 10/24: K30_modVar_bal ... Scenario 11/24: K30_hiVar_bal ... Scenario 12/24: K30_modVar_imbal ... Scenario 13/24: K30_hiVar_imbal ... Scenario 14/24: K30_modVar_int ... Scenario 15/24: K30_hiVar_int ... Scenario 16/24: K30_hiVar_int_imb ... Scenario 17/24: H0_K10_noVar_bal ... Scenario 18/24: H0_K10_modVar_bal ... Scenario 19/24: H0_K10_hiVar_bal ... Scenario 20/24: H0_K10_hiVar_imbal ... Scenario 21/24: H0_K30_noVar_bal ... Scenario 22/24: H0_K30_modVar_bal ... Scenario 23/24: H0_K30_hiVar_bal ... Scenario 24/24: H0_K30_hiVar_imbal ...

Each scenario is replicated $R = 1,500$ times. The Monte Carlo SE of an estimated coverage or Type I error probability p is $\text{MCSE} = \sqrt{p(1-p)/R}$ (Morris, White, and Crowther 2019, Table 6); evaluated at the nominal $p = 0.95$ (worst case among the coverage targets considered here), $R \geq 1,320$ is required for $\text{MCSE} \leq 0.6$ percentage points, and $R = 1,500$ gives $\text{MCSE} \approx 0.56$ pp, satisfying the target with a margin. Treatment is assigned via balanced allocation within each site (equal numbers of treatment and control participants, or within ± 1 when n_i is odd). All random number generation uses a fixed seed (`set.seed(20260310)`) and the L'Ecuyer-CMRG RNG for reproducibility. The `sim-run` chunk's cache is keyed on the MD5 hash of `analysis/scripts/sim_multicenter.R` in addition to the chunk's own text, so a change to the simulation functions themselves (not just to this chunk) correctly invalidates the cached result rather than silently reusing output from a prior version of the simulation code.

3 Results

3.1 Headline findings

We observe that across 16 alternative- hypothesis scenarios ($\delta = 0.30$) and $R = 1,500$ replications per scenario, the random-site mixed model attained power be-

tween 0.54 and 0.97, with coverage of the 95% confidence interval ranging from 0.895 to 0.960 and absolute bias never exceeding 0.011 (Morris Table 6 Monte Carlo SEs are given in parentheses next to the power and coverage point estimates in Table 1). The fixed-site estimator closely tracked the random-site estimator. Ignoring site entirely produced overcoverage (up to 0.984) and correspondingly lower power (as low as 0.40) in scenarios with non-trivial between-site variance. The lowest convergence rate (hard `lmer` errors only) observed across any scenario or method was 1.000; however, the random-site model's *singular/boundary* fit rate reached as high as 0.557 in the low- and no-site-variance scenarios, where the true random-intercept variance is at or near the parameter-space boundary. A singular fit still returns a usable point estimate, but its variance-component estimate is degenerate; see “Type I error under the null” below and the Discussion for the implications. Under the null-hypothesis scenarios ($\delta = 0$; Table 2), the empirical Type I error rate for the random-site method ranged from 0.041 to 0.068, and for the ignore-site method from 0.014 to 0.047, against a nominal $\alpha = 0.05$.

3.2 Summary of performance metrics

3.3 Bias




Figure 1: Bias of the treatment effect estimate across scenarios and analytic methods. Dashed line at zero indicates no bias.

3.4 Power and coverage

3.5 Impact of imbalanced site sizes

3.6 Type I error under the null

Coverage of a 95% CI centered on the true alternative ($\delta = 0.30$, Table 1) is not the same quantity as the Type I error rate of the associated test under a true




Figure 2: Power (left) and coverage probability (right) across scenarios. Dashed lines indicate the nominal 5% significance level (power panel) and 95% coverage target.




Figure 3: Comparison of balanced versus imbalanced site sizes for the random effects method. Panels show empirical standard error (left) and power (right).

Table 1: Simulation results under the alternative hypothesis ($\delta = 0.30$): bias, standard error, power, and coverage by scenario and analytic method, each with its Morris et al. (2019) Table 6 Monte Carlo SE in parentheses where applicable. R = 1500 replications per scenario. ‘Conv.’ is the rate of successful `lmer` fits (errors only); ‘Sing.’ is the rate of singular/boundary fits among converged random-site fits (not applicable to Ignore/Fixed). Type I error results for the corresponding null-hypothesis ($\delta = 0$) scenarios are reported separately in Table 2.

Scenario	Method	Bias	Emp. SE	Model SE	Power (MCSE)	Coverage (MCSE)	Conv.	Sing.
K10 hiVar bal	Fixed	-0.003	0.140	0.141	0.547 (0.013)	0.960 (0.005)	100.0%	0.0%
K10 hiVar bal	Ignore	-0.003	0.140	0.169	0.400 (0.013)	0.984 (0.003)	100.0%	0.0%
K10 hiVar bal	Random	-0.003	0.140	0.141	0.548 (0.013)	0.960 (0.005)	100.0%	0.0%
K10 hiVar imbal	Fixed	-0.001	0.141	0.141	0.558 (0.013)	0.951 (0.006)	100.0%	0.0%
K10 hiVar imbal	Ignore	-0.001	0.142	0.167	0.413 (0.013)	0.979 (0.004)	100.0%	0.0%
K10 hiVar imbal	Random	-0.001	0.141	0.141	0.557 (0.013)	0.953 (0.005)	100.0%	0.1%
K10 hiVar int	Fixed	0.011	0.148	0.142	0.572 (0.013)	0.939 (0.006)	100.0%	0.0%
K10 hiVar int	Ignore	0.011	0.148	0.171	0.428 (0.013)	0.972 (0.004)	100.0%	0.0%
K10 hiVar int	Random	0.011	0.148	0.142	0.573 (0.013)	0.939 (0.006)	100.0%	0.0%
K10 hiVar int imb	Fixed	-0.007	0.173	0.142	0.537 (0.013)	0.901 (0.008)	100.0%	0.0%
K10 hiVar int imb	Ignore	-0.006	0.173	0.168	0.415 (0.013)	0.934 (0.006)	100.0%	0.0%
K10 hiVar int imb	Random	-0.007	0.173	0.142	0.537 (0.013)	0.901 (0.008)	100.0%	0.2%
K10 modVar bal	Fixed	0.002	0.141	0.141	0.561 (0.013)	0.949 (0.006)	100.0%	0.0%
K10 modVar bal	Ignore	0.002	0.141	0.148	0.533 (0.013)	0.955 (0.005)	100.0%	0.0%
K10 modVar bal	Random	0.002	0.141	0.141	0.561 (0.013)	0.949 (0.006)	100.0%	3.2%
K10 modVar imbal	Fixed	-0.002	0.139	0.141	0.553 (0.013)	0.957 (0.005)	100.0%	0.0%
K10 modVar imbal	Ignore	-0.002	0.139	0.147	0.515 (0.013)	0.965 (0.005)	100.0%	0.0%
K10 modVar imbal	Random	-0.002	0.139	0.141	0.550 (0.013)	0.958 (0.005)	100.0%	8.2%
K10 modVar int	Fixed	0.004	0.156	0.142	0.559 (0.013)	0.928 (0.007)	100.0%	0.0%
K10 modVar int	Ignore	0.004	0.156	0.149	0.528 (0.013)	0.942 (0.006)	100.0%	0.0%
K10 modVar int	Random	0.004	0.156	0.142	0.559 (0.013)	0.928 (0.007)	100.0%	3.1%
K10 noVar bal	Fixed	0.003	0.140	0.142	0.569 (0.013)	0.951 (0.006)	100.0%	0.0%
K10 noVar bal	Ignore	0.003	0.140	0.142	0.568 (0.013)	0.950 (0.006)	100.0%	0.0%
K10 noVar bal	Random	0.003	0.140	0.141	0.573 (0.013)	0.950 (0.006)	100.0%	55.7%
K30 hiVar bal	Fixed	0.001	0.082	0.082	0.947 (0.006)	0.937 (0.006)	100.0%	0.0%
K30 hiVar bal	Ignore	0.001	0.082	0.099	0.898 (0.008)	0.980 (0.004)	100.0%	0.0%
K30 hiVar bal	Random	0.001	0.082	0.082	0.947 (0.006)	0.937 (0.006)	100.0%	0.0%
K30 hiVar imbal	Fixed	0.004	0.081	0.082	0.974 (0.004)	0.947 (0.006)	100.0%	0.0%
K30 hiVar imbal	Ignore	0.004	0.081	0.098	0.919 (0.007)	0.975 (0.004)	100.0%	0.0%
K30 hiVar imbal	Random	0.004	0.081	0.082	0.974 (0.004)	0.948 (0.006)	100.0%	0.0%
K30 hiVar int	Fixed	0.003	0.092	0.082	0.937 (0.006)	0.928 (0.007)	100.0%	0.0%
K30 hiVar int	Ignore	0.003	0.092	0.100	0.879 (0.008)	0.972 (0.004)	100.0%	0.0%
K30 hiVar int	Random	0.003	0.092	0.082	0.937 (0.006)	0.928 (0.007)	100.0%	0.0%
K30 hiVar int imb	Fixed	0.003	0.100	0.082	0.925 (0.007)	0.894 (0.008)	100.0%	0.0%
K30 hiVar int imb	Ignore	0.003	0.101	0.099	0.855 (0.009)	0.941 (0.006)	100.0%	0.0%
K30 hiVar int imb	Random	0.003	0.100	0.082	0.924 (0.007)	0.895 (0.008)	100.0%	0.0%
K30 modVar bal	Fixed	0.001	0.079	0.082	0.963 (0.005)	0.954 (0.005)	100.0%	0.0%
K30 modVar bal	Ignore	0.001	0.079	0.086	0.954 (0.005)	0.963 (0.005)	100.0%	0.0%
K30 modVar bal	Random	0.001	0.079	0.082	0.963 (0.005)	0.954 (0.005)	100.0%	0.1%
K30 modVar imbal	Fixed	0.000	0.081	0.082	0.959 (0.005)	0.956 (0.005)	100.0%	0.0%
K30 modVar imbal	Ignore	0.000	0.081	0.085	0.948 (0.006)	0.961 (0.005)	100.0%	0.0%
K30 modVar imbal	Random	0.000	0.081	0.082	0.956 (0.005)	0.954 (0.005)	100.0%	0.1%
K30 modVar int	Fixed	-0.002	0.091	0.082	0.936 (0.006)	0.923 (0.007)	100.0%	0.0%
K30 modVar int	Ignore	-0.002	0.091	0.086	0.918 (0.007)	0.936 (0.006)	100.0%	0.0%
K30 modVar int	Random	-0.002	0.091	0.082	0.936 (0.006)	0.923 (0.007)	100.0%	0.0%
K30 noVar bal	Fixed	0.000	0.082	0.082	0.959 (0.005)	0.943 (0.006)	100.0%	0.0%
K30 noVar bal	Ignore	0.000	0.082	0.082	0.960 (0.005)	0.943 (0.006)	100.0%	0.0%
K30 noVar bal	Random	0.000	0.082	0.081	0.960 (0.005)	0.943 (0.006)	100.0%	52.7%

373 null effect ($\delta = 0$), which is the quantity most directly relevant to the stratified-
374 randomization literature motivating this study^{11,13}. Table 2 reports the empirical
375 rejection rate at $\alpha = 0.05$ for the reduced null-hypothesis scenario subset described
376 in Methods.

377 All three methods maintain Type I error close to the nominal 5% level across the
378 null-hypothesis scenarios, including the high-site-variance and imbalanced condi-
379 tions, and including the scenarios in which the random-site model’s singular-fit
380 rate is highest. This is consistent with the alternative-hypothesis coverage results
381 in Table 1 and provides direct evidence – rather than an inference from coverage
382 of a shifted CI – that none of the three methods inflates the false-positive rate un-

Table 2: Empirical Type I error under the null hypothesis ($\delta = 0$) at nominal $\alpha = 0.05$, by scenario and analytic method, with Morris et al. (2019) Table 6 Monte Carlo SE in parentheses. R = 1500 replications per scenario. ‘Sing.’ is the random-site singular/boundary-fit rate among converged fits.

Scenario	Method	Bias	Type I error (MCSE)	Coverage (MCSE)	Sing.
K10 hiVar bal	Fixed	0.000	0.051 (0.006)	0.949 (0.006)	0.0%
K10 hiVar bal	Ignore	0.000	0.021 (0.004)	0.979 (0.004)	0.0%
K10 hiVar bal	Random	0.000	0.051 (0.006)	0.949 (0.006)	0.1%
K10 hiVar imbal	Fixed	-0.005	0.048 (0.006)	0.952 (0.006)	0.0%
K10 hiVar imbal	Ignore	-0.005	0.019 (0.003)	0.981 (0.003)	0.0%
K10 hiVar imbal	Random	-0.005	0.046 (0.005)	0.954 (0.005)	0.2%
K10 modVar bal	Fixed	0.003	0.057 (0.006)	0.943 (0.006)	0.0%
K10 modVar bal	Ignore	0.003	0.045 (0.005)	0.955 (0.005)	0.0%
K10 modVar bal	Random	0.003	0.057 (0.006)	0.943 (0.006)	3.3%
K10 noVar bal	Fixed	0.001	0.041 (0.005)	0.959 (0.005)	0.0%
K10 noVar bal	Ignore	0.001	0.041 (0.005)	0.959 (0.005)	0.0%
K10 noVar bal	Random	0.001	0.041 (0.005)	0.959 (0.005)	56.9%
K30 hiVar bal	Fixed	0.001	0.048 (0.006)	0.952 (0.006)	0.0%
K30 hiVar bal	Ignore	0.001	0.014 (0.003)	0.986 (0.003)	0.0%
K30 hiVar bal	Random	0.001	0.048 (0.006)	0.952 (0.006)	0.0%
K30 hiVar imbal	Fixed	-0.004	0.065 (0.006)	0.935 (0.006)	0.0%
K30 hiVar imbal	Ignore	-0.004	0.023 (0.004)	0.977 (0.004)	0.0%
K30 hiVar imbal	Random	-0.004	0.068 (0.007)	0.932 (0.007)	0.0%
K30 modVar bal	Fixed	0.000	0.052 (0.006)	0.948 (0.006)	0.0%
K30 modVar bal	Ignore	0.000	0.042 (0.005)	0.958 (0.005)	0.0%
K30 modVar bal	Random	0.000	0.052 (0.006)	0.948 (0.006)	0.0%
K30 noVar bal	Fixed	0.001	0.047 (0.005)	0.953 (0.005)	0.0%
K30 noVar bal	Ignore	0.001	0.047 (0.005)	0.953 (0.005)	0.0%
K30 noVar bal	Random	0.001	0.048 (0.006)	0.952 (0.006)	53.9%

der the conditions simulated here. It does not, by itself, establish that the t -based reference distribution used for the random-site method (Methods, “Analytic methods”) is correctly calibrated in general; the boundary/singular fits flagged in Table 1 and Table 2 return a point estimate and standard error that are numerically valid but whose sampling distribution at the parameter-space boundary is not guaranteed to match the naive t approximation used here, and this simulation’s null-arm results should not be read as validating that approximation outside the specific scenarios simulated.

3.7 Random slopes under treatment-by-site interaction

The random-site method fitted throughout this report (Table 1, Table 2) includes only a random intercept, which is misspecified for the six treatment-by-site interaction scenarios ($\sigma_\gamma^2 = 0.04$), since those scenarios are generated with a genuine random treatment slope, γ_i (Data generating model). Table 3 compares that intercept-only specification against the correctly specified random-slopes method (Analytic methods, Method 4), fit on the same generated data, restricted to those six scenarios.

Across these six scenarios, the random-slopes method’s coverage relative to the (misspecified) random-intercept method’s coverage ranged from +0.020 to +0.039 percentage points, and its singular-fit rate was 33.6 to 65.4%, reflecting the additional variance and covariance parameters the random-slopes model must estimate

Table 3: Random intercept versus random slopes for the six treatment-by-site-interaction scenarios ($\sigma_\gamma^2 = 0.04$), fit on the same generated data. ‘Sing.’ is the singular/boundary-fit rate among converged fits for that method; for Random slope this reflects the joint intercept-slope covariance parameter space, a higher-dimensional boundary than the single-variance-component boundary tracked for Random intercept.

Scenario	Method	Bias	Emp. SE	Power (MCSE)	Coverage (MCSE)	Sing.
K10 hiVar int	Random intercept	0.011	0.148	0.573 (0.013)	0.939 (0.006)	0.0%
K10 hiVar int	Random slope	0.011	0.148	0.456 (0.013)	0.963 (0.005)	49.7%
K10 hiVar int imb	Random intercept	-0.007	0.173	0.537 (0.013)	0.901 (0.008)	0.2%
K10 hiVar int imb	Random slope	-0.007	0.175	0.417 (0.013)	0.939 (0.006)	58.9%
K10 modVar int	Random intercept	0.004	0.156	0.559 (0.013)	0.928 (0.007)	3.1%
K10 modVar int	Random slope	0.004	0.156	0.446 (0.013)	0.956 (0.005)	65.4%
K30 hiVar int	Random intercept	0.003	0.092	0.937 (0.006)	0.928 (0.007)	0.0%
K30 hiVar int	Random slope	0.003	0.092	0.911 (0.007)	0.948 (0.006)	33.6%
K30 hiVar int imb	Random intercept	0.003	0.100	0.924 (0.007)	0.895 (0.008)	0.0%
K30 hiVar int imb	Random slope	0.003	0.101	0.873 (0.009)	0.934 (0.006)	38.0%
K30 modVar int	Random intercept	-0.002	0.091	0.936 (0.006)	0.923 (0.007)	0.0%
K30 modVar int	Random slope	-0.002	0.091	0.889 (0.008)	0.947 (0.006)	47.9%

relative to the random-intercept model. See the Discussion, “Treatment-by-site interaction,” for interpretation.

4 Discussion

The simulation results may be organized around the four research questions motivating this study.

Modeling site effects. Consistent with prior simulation work^{7,8}, we find that the random effects model attained nominal or near- nominal coverage and the highest power across nearly all scenarios. The fixed effects model performed comparably when the number of sites was small ($K = 10$) but lost efficiency with $K = 30$ due to the large number of nuisance parameters. Ignoring site effects produced unbiased point estimates but inflated standard errors when site-level variance was present, leading to conservative inference and reduced power. These findings support the random effects model as a robust default for multicenter RCTs, as recommended in the literature^{1,16}. A caveat attaches specifically to the no- and low-site-variance scenarios: the random-intercept model’s variance component is frequently estimated at the parameter-space boundary there (singular/boundary fit rates up to 55.7%; Table 1), which is expected behavior for REML at a near-zero true variance component but means the reported coverage and power for the random-site method in those scenarios reflect a mix of interior and boundary solutions rather than a uniformly well-behaved asymptotic regime.

A methodological caveat applies to how this recommendation is justified, not to the simulation’s findings themselves: the recommendation rests on a design-based argument that requires no claim about hospitals being a literal random sample of a larger population, which is the more defensible of two justifications the published literature offers for this choice; see Introduction, “Two different reasons statisticians give for treating site as random,” for the full discussion and

429 citations.

430 **Failure to adjust for stratification.** The ‘ignore site’ method effectively repre-
431 sents a failure to adjust for the stratification variable (site) used in randomization.
432 The overcoverage and power loss observed under this method confirm the findings
433 of Kahan and Morris¹¹: when randomization is stratified by center, the analysis
434 must account for center to obtain valid inference. The magnitude of the power
435 penalty depends on the ICC: with no site-level variance, ignoring site is harmless,
436 but as σ_α^2 increases, the penalty becomes substantial.

437 **Treatment-by-site interaction.** Introducing random treatment-by-site inter-
438 action ($\sigma_\gamma^2 = 0.04$) increased the variability of treatment effect estimates across
439 replications for all methods. The random-intercept model, which does not ex-
440 plicitly model this interaction in the fitted specification, nonetheless maintained
441 reasonable coverage in Table 1 because the inflated variability was partially ab-
442 sorbed into the residual. Table 3 (Results, “Random slopes under treatment-by-
443 site interaction”) reports what the correctly specified random-slopes model does
444 differently on the same data: coverage changed by +0.020 to +0.039 percentage
445 points relative to the random-intercept model across the six interaction scenarios,
446 a modest shift that suggests the random-intercept model’s coverage in Table 1 was
447 not substantially compromised by the omitted slope in the conditions simulated
448 here – but the random-slopes model’s singular-fit rate (33.6–65.4%) was substan-
449 tial, reflecting how much harder the additional variance-covariance parameters
450 are to estimate from a single trial’s worth of data, and this trade-off, not coverage
451 alone, is the more important lesson: a random-slopes model is the theoretically
452 correct choice under known treatment heterogeneity, but it is not a free upgrade,
453 and a researcher choosing it should expect and plan for a materially higher rate of
454 boundary/singular fits than the simpler random-intercept model produces. Jones
455 and colleagues and Fedorov and Jones have argued that some degree of treatment-
456 by-center interaction is ubiquitous, and that the power to detect it is typically
457 inadequate^{4,15}. In our view, researchers should plan for its presence rather than
458 test for its absence, and Table 3’s results suggest that planning should weigh the
459 random-slopes model’s estimation difficulty against its theoretical correctness, not
460 assume the more complex model is unambiguously better in a single trial’s finite
461 sample.

462 **Imbalanced site sizes.** The comparison of balanced versus imbalanced site
463 sizes reveals that for the random effects model, moderate imbalance ($CV \approx 1.1$)
464 produces only a modest increase in empirical standard error and a correspondingly
465 small decrease in power. This finding is consistent with Ruvuna and with Vierron
466 and Giraudeau, who showed that the efficiency loss from unequal site sizes is
467 generally small unless the imbalance is extreme^{18,19}. The more concerning scenario
468 is when site size is confounded with site characteristics – a possibility not simulated
469 here but discussed elsewhere^{15,17}. Under informative site sizes, where larger sites
470 systematically differ from smaller sites, all methods may produce biased estimates
471 of the population-average treatment effect.

Limitations. We should note that this simulation is restricted to continuous outcomes analyzed with normal-theory models. Binary and time-to-event outcomes raise additional complications, including separation in fixed effects logistic regression with many small sites^{5,20,23}. The simulation does not incorporate dropout, non-compliance, or informative site sizes. Table 1’s primary comparison of site-modeling strategies still fits only a random intercept for the random-site method, by design, since that is the question those 16 scenarios address; a random-slopes model is now fit and compared against it, but only for the six treatment-by-site-interaction scenarios (Table 3), where it is the correctly specified model. The null-hypothesis (Type I error) arm covers a reduced subset of eight of the sixteen alternative-hypothesis scenarios, omitting the treatment-by-site-interaction conditions; Type I error under a null average effect combined with genuine site-level interaction variance, analyzed by the random-slopes model, is not assessed here and remains a natural extension of the null arm. The random-site and random-slope methods’ p -values and CIs both use an N -minus-fixed-effects degrees of freedom that is not a Kenward-Roger or Satterthwaite correction (Methods, “Analytic methods”); the null-arm results in Table 2 show this approximation was not anti-conservative for the random-site method in the scenarios simulated (the null arm does not cover the random-slope method, per the point above), but that is not a general guarantee, particularly for the boundary/singular fits discussed above and in Table 3.

5 Future research

Several directions warrant further investigation:

1. **Random slopes models (partially addressed).** A random treatment slope ($\delta + \gamma_i$ with γ_i estimated) is now fit and compared against the random-intercept model for the six treatment-by-site-interaction scenarios (Results, “Random slopes under treatment-by-site interaction”; Table 3). What remains is to extend the null-hypothesis arm to cover interaction scenarios analyzed by the random-slopes model, so that Type I error, not only coverage and singular-fit rate, can be assessed for that method, and to extend the K/σ_α^2 /balance factorial grid used for the random-intercept comparison to the random-slopes method as well, rather than the current six-scenario subset.
2. **Non-normal outcomes.** Binary and count outcomes require generalized linear mixed models or GEE, and the relative performance of methods may differ from the continuous case, particularly with sparse data per site^{20,23}.
3. **Informative site sizes.** If site size is associated with site-level treatment effect (e.g., expert centers enroll more patients and also deliver more effective treatment), standard methods may yield biased population-average estimates. Inverse-probability weighting by site or joint modeling of enrollment and outcome merit investigation¹⁷.

- 513 4. **Small-sample corrections.** The current implementation uses a naive N -
514 minus-fixed-effects degrees of freedom for the random-site method's p -values
515 and CIs (Methods, "Analytic methods"), which performed adequately in the
516 null-arm scenarios simulated here (Table 2) but is not a general substitute
517 for a proper small-sample correction. Kenward-Roger or Satterthwaite de-
518 nominator degrees of freedom corrections (available in `lmerTest`) should be
519 substituted and evaluated across the same scenario grid, particularly for the
520 $K = 10$, high- site-variance conditions where few sites most threaten the
521 adequacy of any df approximation.
- 522 5. **Bayesian approaches.** Hierarchical Bayesian models offer partial pooling
523 of site-specific treatment effects, which may outperform frequentist random
524 effects models when the number of sites is very small⁴.
- 525 6. **Sample size planning.** Tools that incorporate the design effect from mul-
526 ticenter structure, including anticipated ICC and site size imbalance, into
527 prospective sample size calculations would improve trial planning^{19,24}.

528 6 Data and code availability

529 All data reported here are synthetic, generated by `analysis/scripts/sim_multicenter.R`;
530 no external or human-participant data are used. The simulation and reporting
531 code, including this document's source, are version-controlled in the repository
532 containing this report. The simulation is fully reproducible from the fixed
533 seed and RNG kind documented in Methods ("Simulation procedure"); the
534 exact scenario grid used to produce every table and figure in this report is
535 defined in the `sim-run` code chunk of `report.Rmd`. A fixed seed and RNG kind
536 alone do not guarantee identical numeric results across package versions, since
537 `lme4`'s optimizer and convergence-warning behavior (directly relevant to the
538 singular/boundary-fit diagnostics reported here) can change between releases.
539 The exact computational environment, including the R and `lme4` versions used
540 to produce this report, is pinned in `renv.lock` at the repository root; running
541 `renv::restore()` from the repository root recreates that environment exactly.

542 7 Conflict of interest

543 The author declares no conflict of interest. No external funding supported this
544 work.

545 8 References

546 8.1 Morris et al. (2019) ADEMP Compliance

547 We audited this simulation study against the reporting standards proposed
548 by Morris, White, and Crowther²¹. The original audit is recorded at
549 `docs/morris-audit-2026-04-17.md`; a status update following the 2026-08
550 revision described below is appended to that file.

551 **Verdict:** Partially compliant (updated).

552 **Resolved since the original 2026-04-17 audit:**

- 553 • Monte Carlo SEs are now computed for every performance estimate in
554 `summarize_simulation()` (`mcse_bias`, `mcse_empirical_se`, `mcse_mse`,
555 `mcse_model_se`, `mcse_power`, `mcse_coverage`, `mcse_convergence`,
556 `mcse_singular_rate`) and are displayed alongside the corresponding point
557 estimates for power and coverage in Table 1 and Table 2.
- 558 • $R = 1500$ is now justified by an explicit Monte Carlo SE derivation (Methods,
559 “Simulation procedure”).
- 560 • Non-convergence is no longer conflated with dropped-and-ignored replications;
561 `n_converged` and `convergence_rate` are reported explicitly, and
562 paired comparisons across methods are preserved because every method’s
563 row is retained (as NA on non-convergence) rather than filtered out.
- 564 • `RNGkind("L'Ecuyer-CMRG")` is pinned immediately before the single
565 `set.seed()` call.
- 566 • Per-replicate `.Random.seed` values are captured by `run_simulation()`,
567 attached as the `rng_states` attribute of its return value, and persisted
568 to `analysis/data/derived_data/rng_states.rds` (keyed by scenario
569 label and replicate index), so a specific failing or singular replication from
570 a completed run can be re-run in isolation via `with_rng_state()`-style
571 `.Random.seed` restoration without re-executing the whole simulation up to
572 that point.

573 **Gaps identified in the 2026-08 revision, not present in the original audit:**

- 574 • Singular/boundary `lmer` fits were, prior to this revision, silently absorbed
575 into “converged” by the warning-suppression logic in `fit_methods()`; this
576 has been fixed by recording `singular` and `conv_warning` as explicit out-
577 comes, but the fix is new as of this revision and has not yet been audited
578 against Morris §5.1 on a wider scenario set than the one simulated here.
- 579 • Coverage previously used a fixed $z = 1.96$ half-width, inconsistent with
580 the t -based p -values used for power; both now use the same per-method
581 degrees of freedom (Methods, “Analytic methods”), but that shared degrees-
582 of-freedom choice is itself a naive approximation for the random-site method,
583 not a Kenward-Roger or Satterthwaite correction (Future research, item 4).
- 584 • The simulation previously provided no null-hypothesis arm and could not
585 support a Type I error claim; a reduced null-arm scenario subset has been
586 added (Results, “Type I error under the null”), but it does not cover the
587 treatment-by-site-interaction conditions (Discussion, “Limitations”).

588 **Still open:**

- 589 • The random-effects model fitted throughout includes only a random inter-
590 cept; Future research item 1 (a random-slopes extension) has not yet been
591 implemented, so the treatment-by-site-interaction scenarios are analyzed by
592 a specification the paper’s own Discussion and Limitations sections describe
593 as misspecified for that generating mechanism.

595 *Rendered on 2026-08-14 at 12:45 PDT.*

596 *Source: ~/prj/res/07-multicenter-rct/multicenterrct/analysis/report/report.Rmd*

- 597 1. International Council for Harmonisation. ICH E9: Statistical principles for clinical trials. Published online 1998. https://database.ich.org/sites/default/files/E9_Guideline.pdf
- 598 2. Localio AR, Berlin JA, Ten Have TR, Kimmel SE. Adjustments for center in multicenter studies: An overview. *Annals of Internal Medicine*. 2001;135(2):112-123. doi:10.7326/0003-4819-135-2-200107170-00012
- 599 3. Senn SJ. Some controversies in planning and analysing multi-centre trials. *Statistics in Medicine*. 1998;17(15-16):1753-1765. doi:10.1002/(SICI)1097-0258(19980815/30)17:15/16<1753::AID-SIM977>3.0.CO;2-X
- 600 4. Jones B, Teather D, Wang J, Lewis JA. A comparison of various estimators of a treatment difference for a multi-centre clinical trial. *Statistics in Medicine*. 1998;17(15-16):1767-1777. doi:10.1002/(SICI)1097-0258(19980815/30)17:15/16<1767::AID-SIM978>3.0.CO;2-H
- 601 5. Pickering RM, Weatherall M. The analysis of continuous outcomes in multi-centre trials with small centre sizes. *Statistics in Medicine*. 2007;26(30):5445-5456. doi:10.1002/sim.3068
- 602 6. Bates D, Mächler M, Bolker B, Walker S. Fitting linear mixed-effects models using `lme4`. *Journal of Statistical Software*. 2015;67(1):1-48. doi:10.18637/jss.v067.i01
- 603 7. Chu R, Thabane L, Ma J, Holbrook A, Pullenayegum E, Devereaux PJ. Comparing methods to estimate treatment effects on a continuous outcome in multicentre randomized controlled trials: A simulation study. *BMC Medical Research Methodology*. 2011;11:21. doi:10.1186/1471-2288-11-21
- 604 8. Kahan BC, Morris TP. Analysis of multicentre trials with continuous outcomes: When and how should we account for centre effects? *Statistics in Medicine*. 2013;32(7):1136-1149. doi:10.1002/sim.5667
- 605 9. Islam S, Bangdiwala SI. Accounting for center-level effects in multicenter randomized controlled trials. *Trials*. 2024;25:390. doi:10.1186/s13063-024-08202-w
- 606 10. Falissard B, Chavance M. L'effet centre dans les Études multicentriques: Fixe ou aléatoire? *Revue d'Épidémiologie et de Santé Publique*. 1996;44(5):455-464.

- 607 11. Kahan BC, Morris TP. Improper analysis of trials randomised using stratified blocks or minimisation. *Statistics in Medicine*. 2012;31(4):328-340. doi:10.1002/sim.4431
- 608 12. Edgar K, Roberts I, Sharples L. Including random centre effects in design, analysis and presentation of multi-centre trials. *Trials*. 2021;22:357. doi:10.1186/s13063-021-05266-w
- 609 13. Kahan BC, Morris TP. Adjusting for multiple prognostic factors in the analysis of randomised trials. *BMC Medical Research Methodology*. 2013;13(1):99. doi:10.1186/1471-2288-13-99
- 610 14. U.S. Food and Drug Administration. Adjusting for covariates in randomized clinical trials for drugs and biological products: Final guidance for industry. Published online 2023. <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/adjusting-covariates-randomized-clinical-trials-drugs-and-biological-products>
- 611 15. Fedorov V, Jones B. Are certain multicenter randomized clinical trial structures misleading clinical and policy decisions? *Contemporary Clinical Trials*. 2005;26(5):518-529. doi:10.1016/j.cct.2005.05.002
- 612 16. Senn SJ, Lewis RJ. Treatment effects in multicenter randomized clinical trials. *JAMA*. 2019;321(12):1211-1212. doi:10.1001/jama.2019.1553
- 613 17. Neuhaus JM, McCulloch CE. Separating between- and within-cluster covariate effects by using conditional and partitioning methods. *Journal of the Royal Statistical Society: Series B*. 2006;68(5):859-872. doi:10.1111/j.1467-9868.2006.00570.x
- 614 18. Ruvuna F. Unequal center sizes, sample size, and power in multicenter clinical trials. *Drug Information Journal*. 2004;38(4):387-394. doi:10.1177/009286150403800409
- 615 19. Vierron E, Giraudeau B. Design effect in multicenter studies: Gain or loss of power? *BMC Medical Research Methodology*. 2009;9:39. doi:10.1186/1471-2288-9-39
- 616 20. Agresti A, Hartzel J. Strategies for comparing treatments on a binary response with multi-centre data. *Statistics in Medicine*. 2000;19(8):1115-1139. doi:10.1002/(SICI)1097-0258(20000430)19:8<1115::AID-SIM408>3.0.CO;2-X

- 617 21. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in Medicine*. 2019;38(11):2074-2102. doi:10.1002/sim.8086
- 618 22. International Council for Harmonisation. ICH E9(R1): Addendum on estimands and sensitivity analysis in clinical trials to the guideline on statistical principles for clinical trials. Published online 2019. https://database.ich.org/sites/default/files/E9-R1_Step4_Guideline_2019_1203.pdf
- 619 23. Kahan BC. Accounting for centre-effects in multicentre trials with a binary outcome – when, why, and how? *BMC Medical Research Methodology*. 2014;14(1):20. doi:10.1186/1471-2288-14-20
- 620 24. Sverdlov O, Ryznik Y, Anisimov V, et al. Selecting a randomization method for a multi-center clinical trial with stochastic recruitment considerations. *BMC Medical Research Methodology*. 2024;24(1):52. doi:10.1186/s12874-023-02131-z